



An Introduction to Peer-to-Peer Networks

Presentation for
MIE456 - Information Systems
Infrastructure II

Vinod Muthusamy
October 30, 2003




Agenda

- Overview of P2P
 - Characteristics
 - Benefits
- Unstructured P2P systems
 - Napster (Centralized)
 - Gnutella (Distributed)
 - Kazaa/Fasttrack (Super-peers)
- Structured P2P systems (DHTs)
 - Chord
 - Pastry
 - CAN
- Conclusions

Client/Server Architecture

- Well known, powerful, reliable server is a data source
- Clients request data from server
- Very successful model
 - WWW (HTTP), FTP, Web services, etc.





Client/Server Limitations

- Scalability is hard to achieve
 - Presents a single point of failure
 - Requires administration
 - Unused resources at the network edge
-
- P2P systems try to address these limitations




P2P Computing*

- P2P computing is the sharing of computer resources and services by direct exchange between systems.
- These resources and services include the exchange of information, processing cycles, cache storage, and disk storage for files.
- P2P computing takes advantage of existing computing power, computer storage and networking connectivity, allowing users to leverage their collective power to the 'benefit' of all.

* From http://www-sop.inria.fr/mistral/personnel/Robin.Groenevelt/Publications/Peer-to-Peer_Introduction_Feb.ppt

P2P Architecture

- All nodes are both clients and servers
 - Provide and consume data
 - Any node can initiate a connection
- No centralized data source
 - "The ultimate form of democracy on the Internet"
 - "The ultimate threat to copy-right protection on the Internet"





P2P Network Characteristics

- Clients are also **servers and routers**
 - Nodes contribute content, storage, memory, CPU
- Nodes are **autonomous** (no administrative authority)
- Network is **dynamic**: nodes enter and leave the network “frequently”
- Nodes **collaborate directly** with each other (not through well-known servers)
- Nodes have widely **varying capabilities**



P2P Benefits

- Efficient use of resources
 - Unused bandwidth, storage, processing power at the edge of the network
- Scalability
 - Consumers of resources also donate resources
 - Aggregate resources grow naturally with utilization
- Reliability
 - Replicas
 - Geographic distribution
 - No single point of failure
- Ease of administration
 - Nodes self organize
 - No need to deploy servers to satisfy demand (c.f. scalability)
 - Built-in fault tolerance, replication, and load balancing



P2P Applications

- Are these P2P systems?
 - File sharing (Napster, Gnutella, Kazaa)
 - Multiplayer games (Unreal Tournament, DOOM)
 - Collaborative applications (ICQ, shared whiteboard)
 - Distributed computation (Seti@home)
 - Ad-hoc networks




Popular P2P Systems

- Napster, Gnutella, Kazaa, Freenet
- Large scale sharing of files.
 - User A makes files (music, video, etc.) on their computer available to others
 - User B connects to the network, searches for files and downloads files directly from user A
- Issues of copyright infringement

Napster

- A way to share music files with others
- Users upload their list of files to Napster server
- You send queries to Napster server for files of interest
 - Keyword search (artist, song, album, bitrate, etc.)
- Napster server replies with IP address of users with matching files
- You connect directly to user A to download file



* Figure from <http://computer.howstuffworks.com/file-sharing.htm>

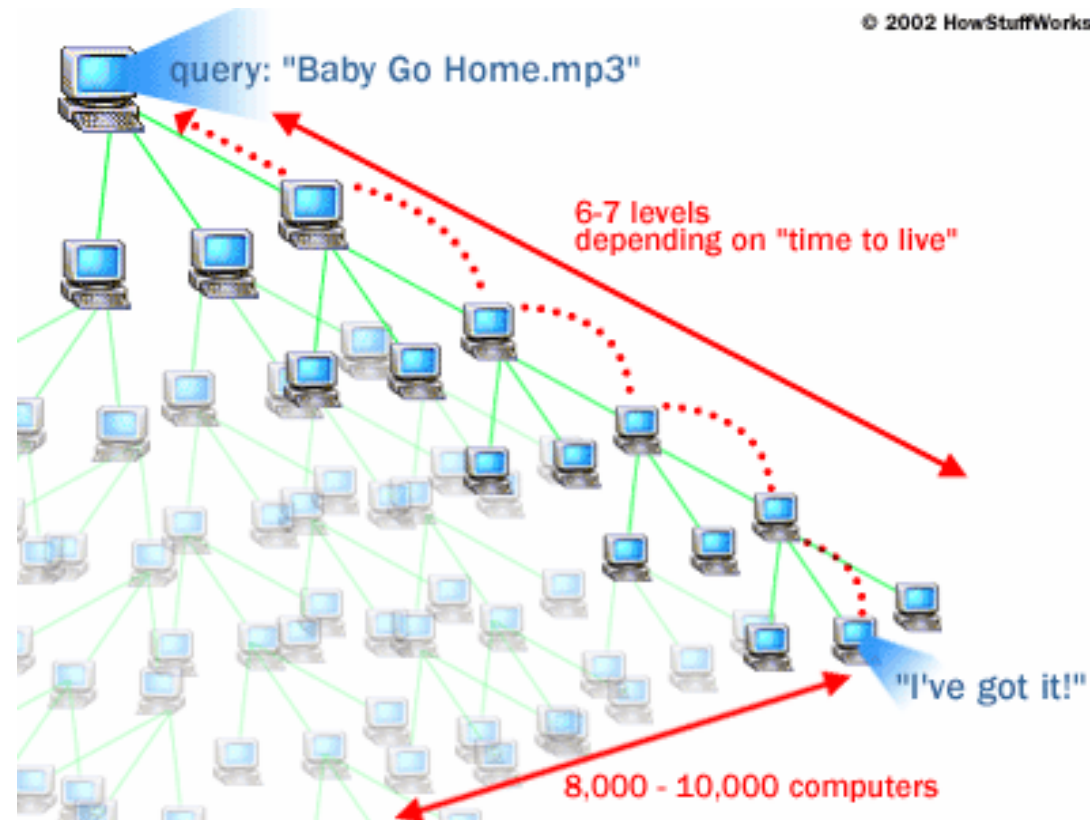
The logo graphic consists of a black crosshair centered on the page. In the top-left quadrant, there is a yellow square. In the bottom-left quadrant, there is a red square. In the bottom-right quadrant, there is a blue square. The word "Napster" is written in a blue, sans-serif font to the right of the crosshair.

Napster

- Central Napster server
 - Can ensure correct results
 - Bottleneck for scalability
 - Single point of failure
 - Susceptible to denial of service
 - Malicious users
 - Lawsuits, legislation
- Search is centralized
- File transfer is direct (peer-to-peer)

Gnutella

- Share any type of files (not just music)
- Decentralized search unlike Napster
- You ask your neighbours for files of interest
- Neighbours ask their neighbours, and so on
 - TTL field quenches messages after a number of hops
- Users with matching files reply to you



* Figure from <http://computer.howstuffworks.com/file-sharing.htm>



Gnutella

- Decentralized
 - No single point of failure
 - Not as susceptible to denial of service
 - Cannot ensure correct results
- Flooding queries
 - Search is now distributed but still not scalable



Kazaa (Fasttrack network)

- Hybrid of centralized Napster and decentralized Gnutella
- Super-peers act as local search hubs
 - Each super-peer is similar to a Napster server for a small portion of the network
 - Super-peers are automatically chosen by the system based on their capacities (storage, bandwidth, etc.) and availability (connection time)
- Users upload their list of files to a super-peer
- Super-peers periodically exchange file lists
- You send queries to a super-peer for files of interest



Free riding*

- File sharing networks rely on users sharing data
- Two types of free riding
 - Downloading but not sharing any data
 - Not sharing any interesting data
- On Gnutella
 - 15% of users contribute 94% of content
 - 63% of users never responded to a query
 - Didn't have "interesting" data

* Data from E. Adar and B.A. Huberman (2000), "Free Riding on Gnutella"



Anonymity

- Napster, Gnutella, Kazaa don't provide anonymity
 - Users know who they are downloading from
 - Others know who sent a query
- Freenet
 - Designed to provide anonymity among other features

- “Smart” queries
 - Requests get routed to correct peer by incremental discovery




Figure 1. Typical request sequence. The request moves through the network from node to node, backing out of a dead-end (step 3) and a loop (step 7) before locating the desired file.



Structured P2P

- Second generation P2P overlay networks
- Self-organizing
- Load balanced
- Fault-tolerant
- Scalable guarantees on numbers of hops to answer a query
 - Major difference with unstructured P2P systems
- Based on a distributed hash table interface



Distributed Hash Tables (DHT)

- Distributed version of a hash table data structure
- Stores (key, value) pairs
 - The key is like a filename
 - The value can be file contents
- Goal: Efficiently insert/lookup/delete (key, value) pairs
- Each peer stores a subset of (key, value) pairs in the system
- Core operation: Find node responsible for a key
 - Map key to node
 - Efficiently route insert/lookup/delete request to this node



DHT Generic Interface

- **Node** id: m-bit identifier (similar to an IP address)
- **Key**: sequence of bytes
- **Value**: sequence of bytes

- put(**key**, **value**)
 - Store (**key,value**) at the **node** responsible for the **key**
- **value** = get(**key**)
 - Retrieve **value** associated with **key** (from the appropriate **node**)



DHT Applications

- Many services can be built on top of a DHT interface
 - File sharing
 - Archival storage
 - Databases
 - Naming, service discovery
 - Chat service
 - Rendezvous-based communication
 - Publish/Subscribe



DHT Desirable Properties

- Keys mapped evenly to all nodes in the network
- Each node maintains information about only a few other nodes
- Messages can be routed to a node efficiently
- Node arrival/departures only affect a few nodes



DHT Routing Protocols

- DHT is a generic **interface**
- There are several **implementations** of this interface
 - Chord [MIT]
 - Pastry [Microsoft Research UK, Rice University]
 - Tapestry [UC Berkeley]
 - Content Addressable Network (CAN) [UC Berkeley]
 - SkipNet [Microsoft Research US, Univ. of Washington]
 - Kademlia [New York University]
 - Viceroy [Israel, UC Berkeley]
 - P-Grid [EPFL Switzerland]
 - Freenet [Ian Clarke]
- These systems are often referred to as **P2P routing substrates** or **P2P overlay networks**




Chord API

- **Node** id: unique m-bit identifier
(hash of IP address or other unique ID)
- **Key**: m-bit identifier (hash of a sequence of bytes)
- **Value**: sequence of bytes

- **API**
 - insert(key, value) → store key/value at r nodes
 - lookup(key)
 - update(key, newval)
 - join(n)
 - leave()


Chord Identifier Circle

- Nodes organized in an **identifier circle** based on node identifiers
- Keys assigned to their **successor** node in the identifier circle
- Hash function ensures even distribution of nodes and keys on the circle




Chord Finger Table

- $O(\log N)$ table size
- i^{th} finger points to first node that succeeds n by at least 2^{i-1}



Chord Key Location

- Lookup in finger table the furthest node that precedes key
- Query homes in on target in $O(\log N)$ hops



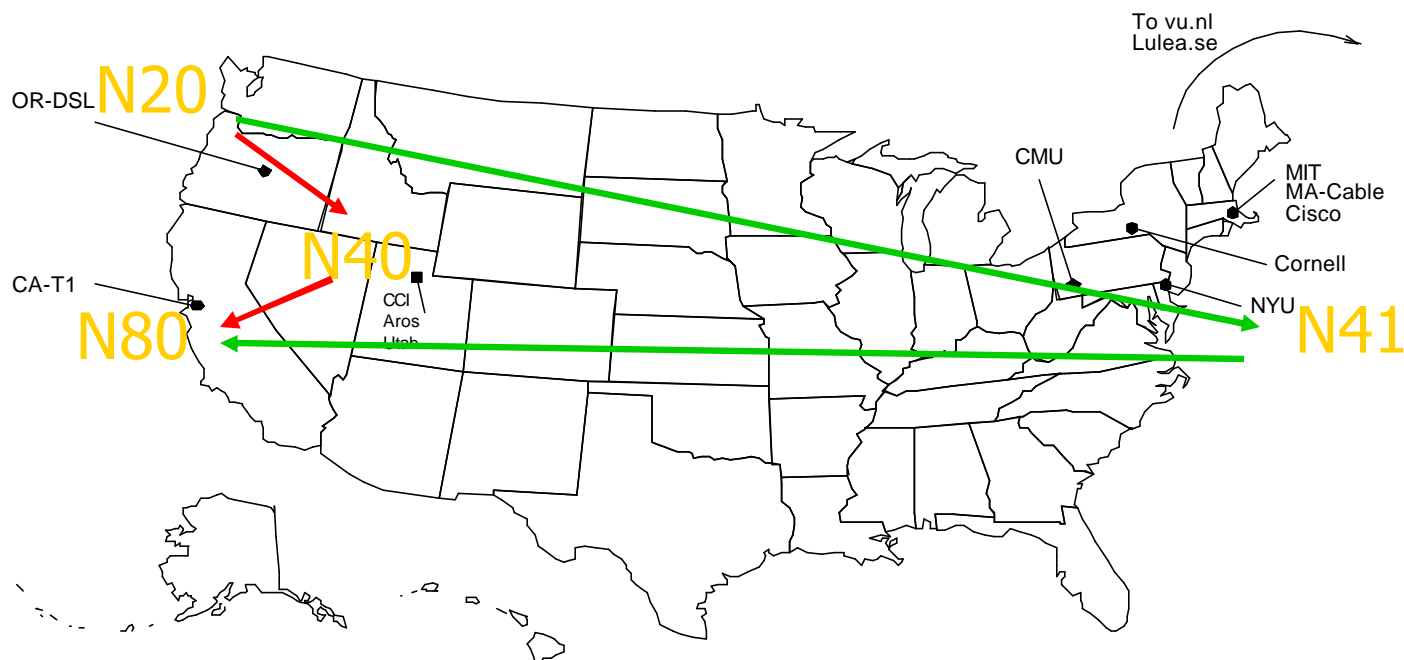


Chord Properties

- In a system with N nodes and K keys, with **high probability**...
 - each node receives at most K/N keys
 - each node maintains info. about $O(\log N)$ other nodes
 - lookups resolved with $O(\log N)$ hops
- No delivery guarantees
- No consistency among replicas
- Hops have poor network locality

Network locality

- Nodes close on ring can be far in the network.



* Figure from <http://project-iris.net/talks/dht-toronto-03.ppt>

Pastry

- Similar interface to Chord
- Considers network **locality** to minimize hops messages travel
- New node needs to know a nearby node to achieve locality
- Each routing hop matches the destination identifier by **one more digit**
 - Many choices in each hop (locality possible)




Figure 2: Routing a message from node $65a1fc$ with key $d46a1c$. The dots depict live nodes in Pastry's circular namespace.

CAN

- Based on a “d-dimensional Cartesian coordinate space on a d-torus”
- Each node owns a distinct **zone** in the space
- Each key hashes to a **point** in the space

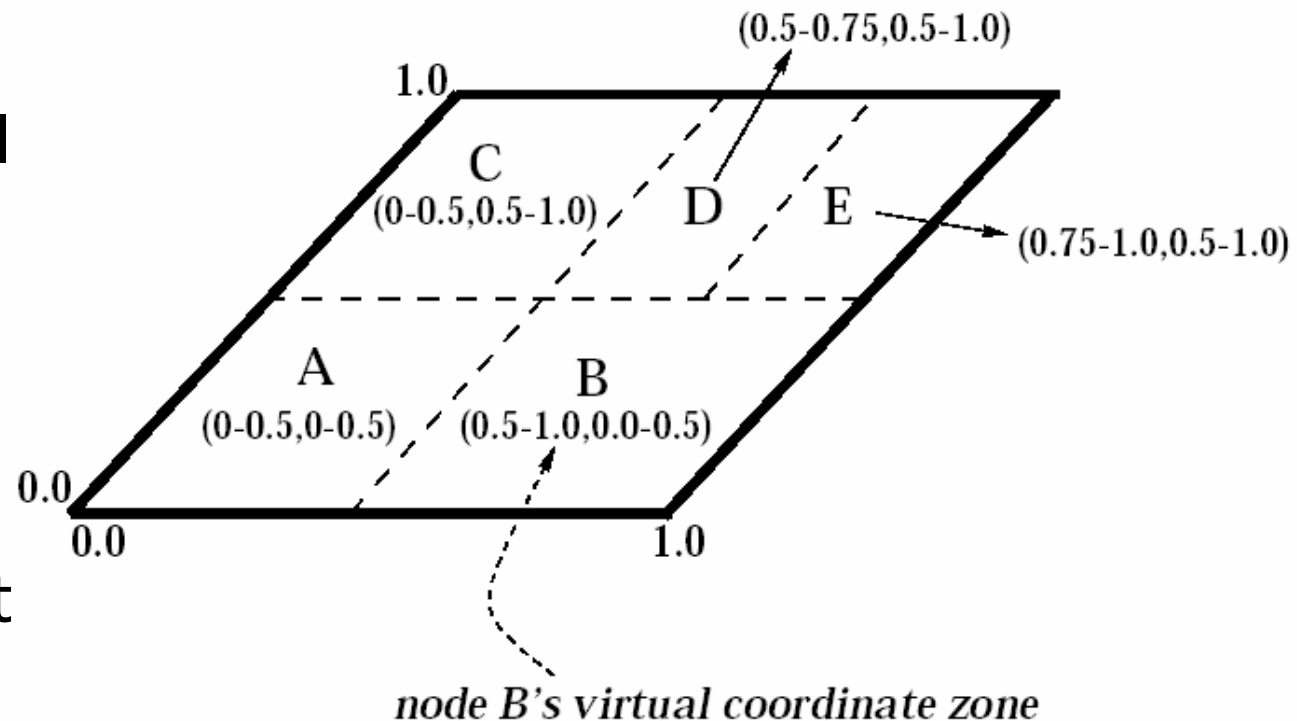




Figure 1: *Example 2-d space with 5 nodes*

CAN Routing and Node Arrival



1's coordinate neighbor set = {2,3,4,5}
 7's coordinate neighbor set = {}

Figure 2: *Example 2-d space before node 7 joins*



1's coordinate neighbor set = {2,3,4,7}
 7's coordinate neighbor set = {1,2,4,5}

Figure 3: *Example 2-d space after node 7 joins*



P2P Review

- Two key functions of P2P systems
 - Sharing content
 - Finding content
- Sharing content
 - Direct transfer between peers
 - All systems do this
 - Structured vs. unstructured placement of data
 - Automatic replication of data
- Finding content
 - Centralized (Napster)
 - Decentralized (Gnutella)
 - Probabilistic guarantees (DHTs)



Conclusions

- P2P connects devices at the edge of the Internet
- Popular in “industry”
 - Napster, Kazaa, etc. allow users to share data
 - Legal issues still to be resolved
- Exciting research in academia
 - DHTs (Chord, Pastry, etc.)
 - Improve properties/performance of overlays
- Applications other than file sharing are being developed